

## Skin Cancer Detection using CLIP

Name : Andi Analta Dwiyanto Palalangan Tunru

ID : 2024962230

The notebook includes the code for both augmented dataset and non-augmented dataset runs. Please check section 2 of the notebook and choose one of the subsections

## 1. Imports

```
!pip install datasets
```

```
Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (3.1.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.16.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.26.4)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (17.0.0)
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.3.8)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.2.2)
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.32.3)
Requirement already satisfied: tqdm>=4.66.3 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.6)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.5.0)
Requirement already satisfied: multiprocessing<0.70.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.70.16)
Requirement already satisfied: fsspec<=2024.9.0,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]<=2024.9.0,
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.11.9)
Requirement already satisfied: huggingface-hub>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.26.3)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (24.2)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.2)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (2.4.4)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (24.2.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.5.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.1.0)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (0.2.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.18.3)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.23.0->
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2.2
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2024
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->datasets)
```

```
from datasets import load_dataset
dataset = load_dataset("marmal88/skin_cancer")
```

```
dataset
```

```
DatasetDict({
  train: Dataset({
    features: ['image', 'image_id', 'lesion_id', 'dx', 'dx_type', 'age', 'sex', 'localization'],
    num_rows: 9577
  })
  validation: Dataset({
    features: ['image', 'image_id', 'lesion_id', 'dx', 'dx_type', 'age', 'sex', 'localization'],
    num_rows: 2492
  })
  test: Dataset({
    features: ['image', 'image_id', 'lesion_id', 'dx', 'dx_type', 'age', 'sex', 'localization'],
    num_rows: 1285
  })
})
```

```
dataset["train"][0]
```

```
{'image': <PIL.JpegImagePlugin.JpegImageFile image mode=RGB size=600x450>,
 'image_id': 'ISIC_0024329',
 'lesion_id': 'HAM_0002954',
 'dx': 'actinic_keratosis',
 'dx_type': 'histo',
 'age': 75.0,
 'sex': 'female',
 'localization': 'lower extremity'}
```

```
def collate_fn(batch):
    pixel_values = torch.stack([example["pixel_values"] for example in batch])
```

```

labels = torch.tensor([example["labels"] for example in batch])
return {"pixel_values": pixel_values, "labels": labels}

```

## 2. Data Preparation

```

import torch
from transformers import CLIPProcessor
from torch.utils.data import DataLoader

processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")

labels = dataset["train"].unique("dx")
label_to_id = {label: idx for idx, label in enumerate(labels)}

processor = CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")

# Set seed value
seed_value = 42

# Seed all random sources
torch.manual_seed(seed_value)
np.random.seed(seed_value)

# For CUDA (if used)
torch.cuda.manual_seed(seed_value)
torch.cuda.manual_seed_all(seed_value) # For multi-GPU
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False

```

### 2.1 Augmented Data

```

from torchvision.transforms import Compose, RandomHorizontalFlip, RandomCrop, ColorJitter, Resize, ToTensor

# Augmentations
train_augmentations = Compose([
    Resize((256, 256)),
    RandomCrop((224, 224)),
    RandomHorizontalFlip(p=0.5),
    ToTensor(),
])

val_augmentations = Compose([
    Resize((224, 224)),
    ToTensor(),
])

from datasets import concatenate_datasets

# Duplicate the training dataset
duplicated_train_dataset = concatenate_datasets([dataset["train"], dataset["train"]])

def preprocess_train(examples):
    num_samples = len(examples["image"])
    half = num_samples // 2

    # Original images (first half)
    original_images = [Resize((224, 224))(img) if not isinstance(img, torch.Tensor) else img for img in examples["image"][:half]]
    original_labels = [label_to_id[label] for label in examples["dx"][:half]]
    original_texts = [f"The lesion shows signs of {label}" for label in examples["dx"][:half]]

    # Augmented images (second half)
    augmented_images = [train_augmentations(img) if not isinstance(img, torch.Tensor) else img for img in examples["image"][half:]]
    augmented_labels = [label_to_id[label] for label in examples["dx"][half:]]
    augmented_texts = [f"The lesion shows signs of {label}" for label in examples["dx"][half:]]

    # Combine images, labels, and texts
    combined_images = original_images + augmented_images
    combined_labels = original_labels + augmented_labels
    combined_texts = original_texts + augmented_texts

    text_inputs = processor(
        text=combined_texts,
        return_tensors="pt",

```

[illegible]



### 3. Training

```

from torch.utils.data import DataLoader

train_dataloader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=32)

from transformers import CLIPModel
from torch import nn, optim
from torch.optim.lr_scheduler import StepLR

model = CLIPModel.from_pretrained("openai/clip-vit-base-patch32")

num_labels = len(labels)
classifier = nn.Linear(model.config.projection_dim, num_labels)

criterion = nn.CrossEntropyLoss()
optimizer = optim.AdamW(model.parameters(), lr=5e-5)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
classifier.to(device)

scheduler = StepLR(optimizer, step_size=5, gamma=0.1)

epochs = 20
for epoch in range(epochs):
    model.train()
    total_loss = 0
    correct_train = 0
    total_train = 0

    for batch in train_dataloader:
        images = batch["pixel_values"].to(device)
        texts = batch["input_ids"].to(device)
        attention_masks = batch["attention_mask"].to(device)
        labels = batch["labels"].to(device)

        outputs = model(pixel_values=images, input_ids=texts, attention_mask=attention_masks, return_dict=True)
        image_embeddings = outputs.image_embeds

        logits = classifier(image_embeddings)

        loss = criterion(logits, labels)

        # Backward pass
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item()

    # Training accuracy
    preds = logits.argmax(dim=-1)
    correct_train += (preds == labels).sum().item()
    total_train += labels.size(0)

# Training metrics
train_accuracy = correct_train / total_train
print(f"Epoch {epoch + 1}/{epochs}, Loss: {total_loss / len(train_dataloader):.4f}, Train Accuracy: {train_accuracy:.4f}, LR: {scheduler.get_last_lr()[0]:.4f}")

scheduler.step()

```



```

Epoch 1/20, Loss: 1.6356, Train Accuracy: 0.6675, LR: [5e-05]
Epoch 2/20, Loss: 1.6263, Train Accuracy: 0.6728, LR: [5e-05]
Epoch 3/20, Loss: 1.6145, Train Accuracy: 0.6863, LR: [5e-05]
Epoch 4/20, Loss: 1.6080, Train Accuracy: 0.6969, LR: [5e-05]
Epoch 5/20, Loss: 1.6037, Train Accuracy: 0.7067, LR: [5e-05]
Epoch 6/20, Loss: 1.5867, Train Accuracy: 0.7343, LR: [5e-06]
Epoch 7/20, Loss: 1.5775, Train Accuracy: 0.7529, LR: [5e-06]
Epoch 8/20, Loss: 1.5747, Train Accuracy: 0.7574, LR: [5e-06]
Epoch 9/20, Loss: 1.5702, Train Accuracy: 0.7646, LR: [5e-06]
Epoch 10/20, Loss: 1.5679, Train Accuracy: 0.7678, LR: [5e-06]
Epoch 11/20, Loss: 1.5625, Train Accuracy: 0.7795, LR: [5.000000000000001e-07]
Epoch 12/20, Loss: 1.5612, Train Accuracy: 0.7824, LR: [5.000000000000001e-07]

```

```
Epoch 13/20, Loss: 1.5608, Train Accuracy: 0.7832, LR: [5.000000000000001e-07]
Epoch 14/20, Loss: 1.5602, Train Accuracy: 0.7844, LR: [5.000000000000001e-07]
Epoch 15/20, Loss: 1.5596, Train Accuracy: 0.7860, LR: [5.000000000000001e-07]
Epoch 16/20, Loss: 1.5595, Train Accuracy: 0.7878, LR: [5.000000000000001e-08]
Epoch 17/20, Loss: 1.5586, Train Accuracy: 0.7894, LR: [5.000000000000001e-08]
Epoch 18/20, Loss: 1.5590, Train Accuracy: 0.7878, LR: [5.000000000000001e-08]
Epoch 19/20, Loss: 1.5581, Train Accuracy: 0.7910, LR: [5.000000000000001e-08]
Epoch 20/20, Loss: 1.5588, Train Accuracy: 0.7889, LR: [5.000000000000001e-08]
```

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
import numpy as np
```

```
# Validation loop with confusion matrix integration
```

```
model.eval()
correct_val = 0
total_val = 0
```

```
all_preds = []
all_labels = []
```

```
for batch in val_dataloader:
    images = batch["pixel_values"].to(device)
    labels = batch["labels"].to(device)
    texts = [f"The lesion shows signs of {label}" for label in batch["labels"]]
```

```
# Preprocess validation text
```

```
text_inputs = processor(
    text=texts,
    return_tensors="pt",
    padding="max_length",
    truncation=True,
    max_length=77,
)
input_ids = text_inputs["input_ids"].to(device)
attention_mask = text_inputs["attention_mask"].to(device)
```

```
with torch.no_grad():
    outputs = model(
        pixel_values=images,
        input_ids=input_ids,
        attention_mask=attention_mask,
        return_dict=True
    )
    image_embeddings = outputs.image_embeds

    logits = classifier(image_embeddings)
    preds = logits.argmax(dim=-1)
```

```
# Collect predictions and labels for confusion matrix
all_preds.extend(preds.cpu().numpy())
all_labels.extend(labels.cpu().numpy())
```

```
# Accuracy calculation
correct_val += (preds == labels).sum().item()
total_val += labels.size(0)
```


```
# Compute validation accuracy
validation_accuracy = correct_val / total_val
print(f"Validation Accuracy: {validation_accuracy:.4f}")
```

```
➡ Validation Accuracy: 0.7917
```

```
# Compute confusion matrix
class_names = ['actinic_keratosis',
'basal_cell_carcinoma',
'benign_keratosis-like_lesions',
'dermatofibroma',
'melanocytic_Nevi',
'melanoma',
'vascular_lesion'] # Replace with actual class names
cm = confusion_matrix(all_labels, all_preds)
```

```
# Plot confusion matrix
plt.figure(figsize=(21, 7))
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
disp.plot(cmap=plt.cm.Blues, values_format="d")
plt.xticks(rotation=45, ha="right")
plt.title("Confusion Matrix")
plt.show()
```

<Figure size 2100x700 with 0 Axes>



```
# Plot class-wise comparison
correct_per_class = np.diag(cm) # Correct predictions per class
total_per_class = cm.sum(axis=1) # Total samples per class
class_accuracy = correct_per_class / total_per_class


plt.figure(figsize=(12, 6))
bar_width = 0.4
x = np.arange(len(class_names))

plt.bar(x - bar_width / 2, total_per_class, width=bar_width, label="Total Samples")
plt.bar(x + bar_width / 2, correct_per_class, width=bar_width, label="Correctly Classified")
plt.xticks(x, class_names)
plt.xticks(rotation=45, ha="right")
plt.ylabel("Number of Samples")
plt.title("Class-wise Comparison of Total vs Correctly Predicted")
plt.legend()
plt.show()


# Plot amount of data for each class in the validation dataset
plt.figure(figsize=(12, 6))
plt.bar(class_names, total_per_class, color='skyblue')
plt.xticks(rotation=45, ha="right")
plt.ylabel("Number of Samples")
plt.title("Number of Samples Per Class in Validation Dataset")
plt.show()
```



Class-wise Comparison of Total vs Correctly Predicted



Number of Samples Per Class in Validation Dataset



```
import matplotlib.pyplot as plt
from PIL import Image

# Filter the dataset based on diagnosis
filtered_dataset = dataset["train"].filter(lambda example: example["dx"] in ["dermatofibroma", "melanoma", "benign_keratosis-like_lesior
```



Filter: 100% 9577/9577 [01:01<00:00, 158.22 examples/s]

```
melano_dataset = filtered_dataset.filter(lambda example: example["dx"] == "melanocytic_Nevi")
```



```

image = melano_dataset[10]["image"] # example["image"] is already a PIL Image

# Resize the image before displaying using Image.resize
image = image.resize((1024, 1024)) # Resize to a smaller size, maintaining aspect ratio

plt.imshow(image)
plt.title(f"Diagnosis: {melano_dataset[0]['dx']}") # Set title with diagnosis, accessing the first element
plt.axis("off") # Hide axes
plt.show()

```



```

dermatofibroma_dataset = filtered_dataset.filter(lambda example: example["dx"] == "dermatofibroma")

image = dermatofibroma_dataset[20]["image"] # example["image"] is already a PIL Image

# Resize the image before displaying using Image.resize
image = image.resize((1024, 1024)) # Resize to a smaller size, maintaining aspect ratio

plt.imshow(image)
plt.title(f"Diagnosis: {dermatofibroma_dataset[0]['dx']}") # Set title with diagnosis, accessing the first element
plt.axis("off") # Hide axes
plt.show()

```



```

melanoma_dataset = filtered_dataset.filter(lambda example: example["dx"] == "melanoma")

image = melanoma_dataset[2]["image"] # example["image"] is already a PIL Image
# Resize the image before displaying using Image.resize
image = image.resize((1024, 1024)) # Resize to a smaller size, maintaining aspect ratio


plt.imshow(image)
plt.title(f"Diagnosis: {melanoma_dataset[0]['dx']}") # Set title with diagnosis, accessing the first element

```

```
plt.axis("off") # Hide axes
```



Diagnosis: melanoma



```
benign_dataset = filtered_dataset.filter(lambda example: example["dx"] == "benign_keratosis-like_lesions")
```

```
image = benign_dataset[5]["image"] # example["image"] is already a PIL Image  
# Resize the image before displaying using Image.resize  
image = image.resize((1024, 1024)) # Resize to a smaller size, maintaining aspect ratio
```


```
plt.imshow(image)  
plt.title(f"Diagnosis: {benign_dataset[0]['dx']}") # Set title with diagnosis, accessing the first element  
plt.axis("off") # Hide axes  
plt.show()
```



Filter: 100%

2234/2234 [00:17<00:00, 126.42 examples/s]

Diagnosis: benign\_keratosis-like\_lesions



- Base Results (Batch Size = 16, Initial Learning Rate = 5e-5)
  - Training Accuracy : 0.835
  - Validation Accuracy : 0.818Indicating that the model isn't experiencing overfitting hence there is no need to implement dropout
- With Regularization (Weight-Decay = 0.01)
  - Training Accuracy : 0.812
  - Validation Accuracy : 0.804